


Two View (Stereo) Reconstruction




Adapted from Szeliski's book

Stereo: epipolar geometry




Fundamental Matrix




- This *epipolar geometry* of two views is described by a very special 3x3 matrix $\mathbf{F}$, called the *fundamental matrix*
- $\mathbf{F}$ maps (homogeneous) *points* in image 1 to *lines* in image 2!
- The epipolar line (in image 2) of point $\mathbf{p}$ is: $\mathbf{Fp}$
- *Epipolar constraint* on corresponding points: $\mathbf{q}^T \mathbf{Fp} = 0$

Fundamental Matrix



- Two Special points: $\mathbf{e}_1$ and $\mathbf{e}_2$ (the *epipoles*): projection of one camera into the other

Fundamental Matrix



- Two Special points: $\mathbf{e}_1$ and $\mathbf{e}_2$ (the *epipoles*): projection of one camera into the other
- All of the epipolar lines in an image pass through the epipole
- Epipoles may or may not be inside the image

Fundamental Matrix

- $\mathbf{F}\mathbf{p}$ is the epipolar line associated with $\mathbf{p}$
- $\mathbf{F}^T\mathbf{q}$ is the epipolar line associated with $\mathbf{q}$
- $\mathbf{F}\mathbf{e}_1 = \mathbf{0}$ and $\mathbf{F}^T\mathbf{e}_2 = \mathbf{0}$
- $\mathbf{F}$ is rank 2
- How many parameters does $\mathbf{F}$ have?


Example – Epipolar lines



Demo


[https://www.cs.cornell.edu/courses/cs5670/2022sp/demos/
FundamentalMatrix/?demo=demo1](https://www.cs.cornell.edu/courses/cs5670/2022sp/demos/FundamentalMatrix/?demo=demo1)

Fundamental Matrix - Derivation



- Why does **F** exist?
- Let's derive it...

Fundamental Matrix - Derivation



$\mathbf{K}_1$: intrinsics of camera 1


$\mathbf{K}_2$: intrinsics of camera 2

$\mathbf{R}$: rotation of image 2 w.r.t. camera 1

$\tilde{p} = \mathbf{K}_1^{-1} p$: ray through p in camera 1's (and world) coordinate system

$\tilde{q} = \mathbf{K}_2^{-1} q$: ray through q in camera 2's coordinate system


Fundamental Matrix - Derivation



- $\tilde{p}$, $R^T \tilde{q}$, and t are coplanar
- epipolar plane can be represented as with its normal $t \times \tilde{p}$

$$(R^T \tilde{q})^T (t \times \tilde{p}) = 0$$

Fundamental Matrix - Derivation



$$(\mathbf{R}^T \tilde{\mathbf{q}})^T (\mathbf{t} \times \tilde{\mathbf{p}}) = 0$$



$$\tilde{\mathbf{q}}^T \mathbf{R} (\mathbf{t} \times \tilde{\mathbf{p}}) = 0$$


Cross-product as linear operator

Useful fact: Cross product with a vector $\mathbf{t}$ can be represented as multiplication with a (*skew-symmetric*) 3x3 matrix

$$[\mathbf{t}]_{\times} = \begin{bmatrix} 0 & -t_z & t_y \\ t_z & 0 & -t_x \\ -t_y & t_x & 0 \end{bmatrix}$$

$$\mathbf{t} \times \tilde{\mathbf{p}} = [\mathbf{t}]_{\times} \tilde{\mathbf{p}}$$


Fundamental Matrix - Derivation



- One more substitution:
 - Cross product with $\mathbf{t}$ (on left) can be represented as a 3x3 matrix

$$[\mathbf{t}]_{\times} = \begin{bmatrix} 0 & -t_z & t_y \\ t_z & 0 & -t_x \\ -t_y & t_x & 0 \end{bmatrix} \quad \mathbf{t} \times \tilde{\mathbf{p}} = [\mathbf{t}]_{\times} \tilde{\mathbf{p}}$$

Fundamental Matrix – Calibrated Cameras




$$\tilde{q}^T \mathbf{R}(t \times \tilde{p}) = 0$$



$$\tilde{q}^T \mathbf{R} [t]_{\times} \tilde{p} = 0$$

Fundamental Matrix – Calibrated Cameras




$\tilde{\mathbf{p}} = \mathbf{K}_1^{-1} \mathbf{p}$: ray through $\mathbf{p}$ in camera 1's (and world) coordinate system

$\tilde{\mathbf{q}} = \mathbf{K}_2^{-1} \mathbf{q}$: ray through $\mathbf{q}$ in camera 2's coordinate system

$$\underbrace{\tilde{\mathbf{q}}^T \mathbf{R} [\mathbf{t}]_{\times}}_{\mathbf{E}} \tilde{\mathbf{p}} = 0 \quad \tilde{\mathbf{q}}^T \mathbf{E} \tilde{\mathbf{p}} = 0$$

$\mathbf{E}$ ← the Essential matrix

Fundamental Matrix – Uncalibrated Cameras



$\mathbf{K}_1$: intrinsics of camera 1


$\mathbf{K}_2$: intrinsics of camera 2

$\mathbf{R}$: rotation of image 2 w.r.t. camera 1

$$\mathbf{q}^T \underbrace{\mathbf{K}_2^{-T} \mathbf{R} [\mathbf{t}]_{\times} \mathbf{K}_1^{-1}}_{\mathbf{F}} \mathbf{p} = 0$$

$\mathbf{F}$ ← the Fundamental matrix


Rectified case



$$\mathbf{R} = \mathbf{I}_{3 \times 3}$$
$$\mathbf{t} = \begin{bmatrix} 1 & 0 & 0 \end{bmatrix}^T$$
$$\mathbf{E} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{bmatrix}$$

Rectification

- reproject image planes onto a common plane
 - plane parallel to the line between optical centers
- pixel motion is horizontal after this transformation
- two homographies, one for each input image reprojection





Original stereo pair



After rectification

Estimating $\mathbf{F}$



- If we don't know $\mathbf{K}_1$, $\mathbf{K}_2$, $\mathbf{R}$, or $\mathbf{t}$, can we estimate $\mathbf{F}$ for two images?
- Yes, given enough correspondences

Estimating $\mathbf{F}$

- The fundamental matrix $\mathbf{F}$ is defined by

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

for any pair of matches $\mathbf{x}$ and $\mathbf{x}'$ in two images.

- Let $\mathbf{x} = (u, v, 1)^T$ and $\mathbf{x}' = (u', v', 1)^T$, $\mathbf{F} = \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix}$
each match gives a linear equation

$$uu'f_{11} + vu'f_{12} + u'f_{13} + uv'f_{21} + vv'f_{22} + v'f_{23} + uf_{31} + vf_{32} + f_{33} = 0$$

Estimating F – 8 point algorithm

$$\begin{bmatrix} u_1 u_1' & v_1 u_1' & u_1' & u_1 v_1' & v_1 v_1' & v_1' & u_1 & v_1 & 1 \\ u_2 u_2' & v_2 u_2' & u_2' & u_2 v_2' & v_2 v_2' & v_2' & u_2 & v_2 & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ u_n u_n' & v_n u_n' & u_n' & u_n v_n' & v_n v_n' & v_n' & u_n & v_n & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix} = 0$$

- Like with homographies, instead of solving $\mathbf{A}\mathbf{f} = 0$, we seek unit length $\mathbf{f}$ to minimize $\|\mathbf{A}\mathbf{f}\|$: least eigenvector of $\mathbf{A}^T \mathbf{A}$.

Estimating $\mathbf{F}$ – 8 point algorithm

- $\mathbf{F}$ should have rank 2
- To enforce that $\mathbf{F}$ is of rank 2, $\mathbf{F}$ is replaced by $\mathbf{F}'$ that minimizes $\|\mathbf{F} - \mathbf{F}'\|$ subject to the rank constraint.
- This is achieved by SVD. Let $\mathbf{F} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$, where

$$\mathbf{\Sigma} = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{bmatrix}, \quad \text{let } \mathbf{\Sigma}' = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

then $\mathbf{F}' = \mathbf{U}\mathbf{\Sigma}'\mathbf{V}^T$ is the solution (closest rank-2 matrix to $\mathbf{F}$)

Estimating F – 8 point algorithm

```
% Build the constraint matrix
A = [x2(1,:)'.*x1(1,:) '   x2(1,:)'.*x1(2,:) '   x2(1,:) '   ...
      x2(2,:)'.*x1(1,:) '   x2(2,:)'.*x1(2,:) '   x2(2,:) '   ...
      x1(1,:) '             x1(2,:) '             ones(npts,1) ];

[U,D,V] = svd(A);

% Extract fundamental matrix from the column of V
% corresponding to the smallest singular value.
F = reshape(V(:,9),3,3)';

% Enforce rank2 constraint
[U,D,V] = svd(F);
F = U*diag([D(1,1) D(2,2) 0])*V';
```

Estimating F – 8 point algorithm

- Pros: linear, easy to implement and fast
- Cons: susceptible to noise

Practical Considerations in Estimating Fundamental Matrix

Problem with 8-point algorithm

$$\begin{bmatrix}
 u_1 u_1' & v_1 u_1' & u_1' & u_1 v_1' & v_1 v_1' & v_1' & u_1 & v_1 & 1 \\
 u_2 u_2' & v_2 u_2' & u_2' & u_2 v_2' & v_2 v_2' & v_2' & u_2 & v_2 & 1 \\
 \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\
 u_n u_n' & v_n u_n' & u_n' & u_n v_n' & v_n v_n' & v_n' & u_n & v_n & 1
 \end{bmatrix}
 \begin{bmatrix}
 f_{11} \\
 f_{12} \\
 f_{13} \\
 f_{21} \\
 f_{22} \\
 f_{23} \\
 f_{31} \\
 f_{32} \\
 f_{33}
 \end{bmatrix} = 0$$

~ 10000 ~ 10000 ~ 100 ~ 10000 ~ 10000 ~ 100 ~ 100 ~ 100 1




Orders of magnitude difference
 between column of data matrix
 → least-squares yields poor results

Normalized 8-point algorithm

normalized least squares yields good results

Transform image to $\sim [-1,1] \times [-1,1]$



Normalized 8-point algorithm

- Transform input by $\hat{\mathbf{x}}_i = \mathbf{T}\mathbf{x}_i$ $\hat{\mathbf{x}}'_i = \mathbf{T}\mathbf{x}'_i$
- Call 8-point on $\hat{\mathbf{x}}_i, \hat{\mathbf{x}}'_i$ to obtain $\hat{\mathbf{F}}$
- $\mathbf{F} = \mathbf{T}'^T \hat{\mathbf{F}} \mathbf{T}$

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$
$$\hat{\mathbf{x}}'^T \mathbf{T}'^{-T} \mathbf{F} \mathbf{T}^{-1} \hat{\mathbf{x}} = 0$$

$\hat{\mathbf{F}}$

Normalized 8-point algorithm

```
[x1, T1] = normalise2dpts(x1);  
[x2, T2] = normalise2dpts(x2);
```

```
A = [x2(1,:)'.*x1(1,:) '   x2(1,:)'.*x1(2,:) '   x2(1,:) '   ...  
      x2(2,:)'.*x1(1,:) '   x2(2,:)'.*x1(2,:) '   x2(2,:) '   ...  
      x1(1,:) '             x1(2,:) '             ones(npts,1)  ];
```

```
[U,D,V] = svd(A);
```

```
F = reshape(V(:,9),3,3)';
```

```
[U,D,V] = svd(F);
```


```
F = U*diag([D(1,1) D(2,2) 0])*V';
```

```
% Denormalise
```

```
F = T2'*F*T1;
```



Results (ground truth)

■ Ground truth with standard stereo calibration





Results (ground truth)

■ 8-point algorithm



Results (normalized 8-point algorithm)

■ Normalized 8-point algorithm



What about more than two views?

- The geometry of three views is described by a $3 \times 3 \times 3$ tensor called the *trifocal tensor*
- The geometry of four views is described by a $3 \times 3 \times 3 \times 3$ tensor called the *quadrifocal tensor*
- After this it starts to get complicated...