




CS 229: Machine Learning

Spring quarter – March 30,
2026

Instructor: Greg {Ver Steeg

TA:

Partha Thakuria



Karpathy:
“LLM research is [about] summoning
ghosts”

Amodei (Anthropic CEO):
“Imagine a country of geniuses in a
datacenter”

Just next-token prediction!

And we are *just* a bundle of neurons

What makes it work?

- Powerful enough model
- Trained with large and complex enough data
- Optimizing right thing in right way, and
- Engineered to scale



First week or two – getting into some details

Some other topics:

NLP

- Tuning (LoRA) vs in context learning
- Scaling laws

Vision

self-supervision/contrastive


Uncertainty quantification

Deep Bayes

Optimization dynamics

Gen. models (focus on flow matching)

Graph neural nets



Course logistics

On course website

<https://elearn.ucr.edu/courses/228278>

Syllabus is there

Announcement via Canvas email list

(You should have gotten one already! If not email me to be added.)

3.5 Homeworks

Homework 0	– 10%	<i>(Discuss shortly)</i>
Homework 1	– 15%	(CalcGPT)
Homework 2	– 15%	(Uncertainty quantification)
Homework 3	– 15%	(Flow matching w/ ViT)
	55%	

HW 1-3 will have some significant but challenging Extra Credits for people who want to go deep.

[Restrictions apply](#)

Pay As You Go

\$9.99 for 100 Compute Units

\$49.99 for 500 Compute Units

You currently have 0 compute units.

Compute units expire after 90 days.
Purchase more as you need them.

- ✓ No subscription required.
Only pay for what you use.
- ✓ Faster GPUs
Upgrade to more powerful GPUs.

Recommended

Colab Pro

\$9.99 per month


No cost for students and educators

- ✓ 100 compute units per month
Compute units expire after 90 days.
Purchase more as you need them.
- ✓ Faster GPUs
Upgrade to more powerful GPUs.
- ✓ More memory
Access our highest memory machines.

Free Colab
Pro for
students!
It has
serious
GPUs!

Quizzes – don't study (20%)

- Frequent low stakes feedback for you and me
- Do it without help – value of being wrong
- Practice for final
- New ideas
- *** Challenge problems



Final (20%)

Variations of quiz questions

GPT-4 Prompt: I need a final version of the anthropomorphized paper clip. This time he's taking a final exam and his stress levels are through the roof - ten times the anxiety of the previous pictures.



Generative AI policy / thoughts

- **Yes:** Explain
- **No:** Auto-generate

Experimental code comprehension test (5%)

Background

Quiz 0: Due Saturday. I'll overview results on Monday.
Meant to gauge background knowledge

HW 0: due April 10 (Friday, in 12 days!)

*If you feel like you are not prepared for the class,
consider auditing and letting someone on the waitlist
join.*

Canvas- Modules

Video coding demos

Below are some links to the coding demos I did in CS 224. The actual code is in the "demos" dropbox folder also linked in modules.

You may have to skip the first part of the lecture, where I usually talked about quiz results, to get to the coding part.

Note: they are in *reverse* order, so the bottom video is week 1.


Week 1 - start at 61:35 - tensor basics in pytorch

Week 2 - start at 34:50 - models in pytorch, tokenizer

Week 3 - start at 31:20 - Cross validation (skippable)

week 4 - start at 9:31 - Multilayer perceptron, Loss, SGD (Start here if you have some experience)

week 5 - start at 17:55 - Debugging neural nets, MLP training demo (Start here if you have lots of experience and want a refresher)



Slides

Posted from 1-100 minutes before class

Link is on canvas

Books and papers

Books with good background on various topics online – we won't closely follow one

I'll email some papers to “read” that I'll cover in class



Nano banana

Office hours – see syllabus

Office hours:

1:50pm-3:00pm Monday (right after class - we can walk to my office MRB 4114 for longer chats)

6:00-7:00 pm Thursday on Zoom

TA Office hours:

10:00-11:00 am Tuesday, WCH 110

6:00-7:00 pm Friday on Zoom



HW 0: AI generated name detector


- Catch up on basic Pytorch if you are new to it
- Dive into some LLM details (tokenization, padding, batch processing) for the rest of us

Data comes from Karpathy's "makemore" project

Social security database:
eriksen,anara,lucie...

Generated:
jieron,slan,keyanvexy...

Also check out his
"nanoGPT"
"nanoChat"
"Autoresearch"!



The spelled-out intro to language modeling: building makemore

Andrej Karpathy

18K

Share

Ask

Save

Download

...

github.com/karpathy/makemore




makemore

Public


Watch 38

How to process these? jieron, slan, keyanvexy ...



General problem - turn everything into numbers, turn numbers back into everything

Multimodal LLMs



Turn text into a sequence of numbers

Keyanvexy


11, 5, 25, 1, 14, 22, 5, 24, 25

Then embed as *float vectors*

$E_{11}, E_5, E_{25}, \dots$

Size of the final representation?
(length of sequence, d)

d - embed dimension



Row for
each
token in
vocabula
ry

E, Embedding table

Multiple sequences in a tensor?

keyanvexy

11, 5, 25, 1, 14, 22, 5, 24, 25

greg

7, 18, 5, 7

greg

7, 18, 5, 7, 0, 0, 0, 0, 0

Pad to same length. Pad_token_id=0


```
input_ids = torch.tensor([11, 5, 25, 1, 14, 22, 5, 24, 25],  
                          [ 7, 18, 5, 7, 0, 0, 0, 0, 0])
```

Then embed

Size of the final representation?

(Batch size, max length of sequence,

d – embed dimension



Row for
each
token in
vocabula
ry

E, Embedding table

Best way to tokenize?

Byte-Pair Encoding (BPE)

character-based
models

vocabulary size

word-based
models

f a s t e r
6 1 19 20 5 18

fast er
19731 288

faster
18277

f a s t e s t
6 1 19 20 5 19 20

fast est
19731 791

fastest
1729

q u i c k e s t
17 21 9 3 11 5 19 20


quick est
1550 791

quickest
65536


“keyanvexy”
sequence length
is 9!

No
“keyanvexy”
in

Zooming in on GPT models




High level picture of a transformer block



18 27 11
An -ar -a
Anara

- # parameters does not grow with sequence length
- Can process arbitrary length sequences, in principle
- Pad to get constant length during training


Better picture



- Most parts are processed independently, in parallel EXCEPT Self-attention
 - # parameters does not grow with sequence length
 - Can process arbitrary length sequences, in principle
 - (None of these operations know order)
- Embedding dimension
Row for each token in vocabulary
E, Embedding table

Self-attention

First, split up embedding into three parts



Output depends on a probability distribution over the inputs


$$o_i = \sum_j p(j|i) v_j$$

How do we get the probability distribution? We always use *softmax* to get discrete probabilities


$$p(j|i) = \text{softmax}_j(q_i \cdot k_j) = \frac{e^{q_i \cdot k_j}}{\sum_{j'} e^{q_i \cdot k_{j'}}}$$

Position / order is not used

Permuting inputs would permute outputs



Position



- Most parts are processed independently, in parallel EXCEPT Self-attention
 - # parameters does not grow with sequence length
 - Can process arbitrary length sequences, in parallel
- Example
- Row for each token in vocabulary
For each position in the sequence

Sinusoidal Positional Encoding

Sinusoidal positional encoding is defined as follows, given the token position $i = 1, \dots, L$ and the dimension $\delta = 1, \dots, d$:

$$\text{PE}(i, \delta) = \begin{cases} \sin\left(\frac{i}{10000^{2\delta'/d}}\right) & \text{if } \delta = 2\delta' \\ \cos\left(\frac{i}{10000^{2\delta'/d}}\right) & \text{if } \delta = 2\delta' + 1 \end{cases}$$





Fig. 3. Sinusoidal positional encoding with $L = 32$ and $d = 128$. The value is between -1 (black) and 1 (white) and the value 0 is in gray.

<https://lilianweng.github.io/posts/2023-01-27-the-transformer-family-v2/>

Newer work uses “rotation based” position embedding:

Rotation based position embedding?



Continue down the GPT
rabbit hole next time