


CS-229

Self supervised vision
nets intro

Confidence in neural



Inductive biases in vision: invariance and equivariance

Invariance

- A function $f[x]$ is **invariant** to a transformation $t[]$ if:

$$f[t[x]] = f[x]$$

i.e., the function output is the same even after the transformation is applied.

Invariance example

e.g., Image classification

- Image has been translated, but we want our classifier to give the same result



Equivariance

- A function $f[x]$ is **equivariant** to a transformation $t[]$ if:

$$f[t[x]] = t[f[x]]$$

i.e., the output is transformed in the same way as the input




Equivariance example

e.g., Image segmentation

- Image has been translated and we want segmentation to translate with it





Semi-supervised Vision



How to exploit the large number of unlabeled examples given in semi-supervised learning?

Semi-supervised learning with pseudo labeling and augmenting





Data augmentation in computer vision

<https://pytorch.org/vision/stable/transforms.html>

Very useful in self/semi-supervision, but a few problems:

- Hard to replicate this success in non-vision domains
- We should be a little careful, because some augmentation could change the label.

Original image



Original image



Original image




FixMatch: Simplifying Semi-Supervised Learning with Consistency and Confidence

Kihyuk Sohn* David Berthelot* Chun-Liang Li Zizhao Zhang Nicholas Carlini

Ekin D. Cubuk Alex Kurakin Han Zhang Colin Raffel

Google Research

{kihyuks, dberth, chunliang, zizhaoz, ncarlini,
cubuk, kurakin, zhanghan, craffel}@google.com



MixMatch: A Holistic Approach to Semi-Supervised Learning

David Berthelot
Google Research
dberth@google.com

Nicholas Carlini
Google Research
ncarlini@google.com

Ian Goodfellow
Work done at Google
ian-academic@mailfence.com


Avital Oliver
Google Research

Nicolas Papernot
Google Research

Colin Raffel
Google Research

Loss = cross entropy (on labeled examples) +

$$\mathcal{L}_u = \frac{1}{L|\mathcal{U}'|} \sum_{u, q \in \mathcal{U}'} \|q - p_{\text{model}}(y | u; \theta)\|_2^2$$




1. *Consistency regularization*: Encourage the model to output the same predictions on perturbed unlabeled samples.
2. *Entropy minimization*: Encourage the model to output confident predictions on unlabeled data.
3. *MixUp* augmentation: (coming up)


mixup: BEYOND EMPIRICAL RISK MINIMIZATION

Hongyi Zhang
MIT

Moustapha Cisse, Yann N. Dauphin, David Lopez-Paz*
FAIR




$$\text{mixup}_\lambda(\mathbf{x}_i, \mathbf{x}_j) = \lambda \mathbf{x}_i + (1 - \lambda) \mathbf{x}_j$$
$$p(\text{mixup}_\lambda(y \mid \mathbf{x}_i, \mathbf{x}_j)) \approx \lambda p(y \mid \mathbf{x}_i) + (1 - \lambda) p(y \mid \mathbf{x}_j)$$



mixup: BEYOND EMPIRICAL RISK MINIMIZATION

Hongyi Zhang
MIT

Moustapha Cisse, Yann N. Dauphin, David Lopez-Paz*
FAIR



My interpretation – mixup constrains the shape of the classifier between these manifolds


Improves adversary robustness and confidence calibration too (Wednesday?)

Self-supervised Computer Vision


Maybe you can't pretrain on a similar large labeled dataset.
Then you should consider pretraining on a large *self-supervised* task

We know what to do for LLMS (next word prediction), what do we do in vision?

Why/when does learning a self-supervised embedding help?




Self-supervised pipeline evaluation



Self supervised, or "pre-text" classification task

(1) Self-supervised training



$$\hat{y} = W z + b$$


Train linear model using z , to predict target task

(2) Transfer to target task

This is how self-supervised methods are evaluated in literature


History of things people
tried

Predict rotations!



**Unsupervised Representation Learning by
Predicting Image Rotations**

Spyros Gidaris, Praveer Singh, Nikos Komodakis



$$X = (\text{cat_face}, \text{cat_eye}); Y = 3$$

Predict context

Example:



Question 1:




Question 2:



Unsupervised Visual Representation Learning by Context Prediction

Carl Doersch, Abhinav Gupta, Alexei A. Efros


Solve puzzles



Permutation Set

index	permutation
64	9,4,6,8,3,2,5,1,7


Reorder patches according to the selected permutation



Unsupervised Learning of Visual Representations by Solving Jigsaw Puzzles


Mehdi Noroozi, [Paolo Favaro](#)

Contrastive learning – triplet loss



$$\mathcal{L}_{\text{triplet}}(\mathbf{x}, \mathbf{x}^+, \mathbf{x}^-) = \sum_{\mathbf{x} \in \mathcal{X}} \max(0, \|f(\mathbf{x}) - f(\mathbf{x}^+)\|_2^2 - \|f(\mathbf{x}) - f(\mathbf{x}^-)\|_2^2 + \epsilon)$$

Make these close Make these far apart



To do it in a more self-supervised way, the positive could be a modification of the anchor image, and the negative could be


More popular/modern form of contrastive loss: InfoNCE / CPC

“Info Noise Contrastive Estimation”, “Contrastive Predictive Coding”

Original paper: <https://arxiv.org/abs/1807.03748>

Predict a match in a batch with one positive and many negatives

Using some neural network critic function, f , a classifier



$$\mathcal{L}_{\text{InfoNCE}} = -\mathbb{E} \left[\log \frac{f(\mathbf{x}^+; \mathbf{c})}{\sum_{\mathbf{x}' \in X} f(\mathbf{x}', \mathbf{c})} \right]$$

(Often something like cosine similarity between
embeddings)

Contrastive loss math

We want to define a probability distribution – predict which of the images on the right is correctly matched with left.

First, use a similarity measure. Dot product is common:

$$\text{Sim}(a, i) = z_a \cdot z_j$$

How do we turn it into a distribution?

$$p(y = j) = \text{softmax sim}(z_a, z_j) = \frac{e^{z_a \cdot z_j}}{\sum_i e^{z_a \cdot z_i}}$$




Now, we have to compare this to the *true* answer.

We have set it in this case, $y=1$, or $q(y=1) = 1$,

$q(y \neq 1) = 0$.


Cross entropy:

$$H(p, q) = -\sum_y q(y) \log p(y) = -\log \frac{e^{z_a \cdot z_1}}{\sum_i e^{z_a \cdot z_i}}$$



SimCLR – Self-supervised contrastive similarity learning

SimCLR Framework



Using an InfoNCE style loss

cosine similarity). Then the loss function for a positive pair of examples (i, j) is defined as

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(z_i, z_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(z_i, z_k)/\tau)}, \quad (1)$$

SimCLR augmentations



(a) Original




(b) Crop and resize



(c) Crop, resize (and flip)



(d) Color distort. (drop)



(e) Color distort. (jitter)



(f) Rotate $\{90^\circ, 180^\circ, 270^\circ\}$




(g) Cutout



(h) Gaussian noise



(i) Gaussian blur



(j) Sobel filtering

SimCLRv2

Big Self-Supervised Models are Strong Semi-Supervised Learners

Ting Chen, Simon Kornblith, Kevin Swersky, Mohammad Norouzi, Geoffrey Hinton
Google Research, Brain Team





Figure 2: Top-1 accuracy of previous state-of-the-art (SOTA) methods [1, 2] and our method (SimCLRv2) on ImageNet using only 1% or 10% of the labels. Dashed line denotes fully supervised ResNet-50 trained with 100% of labels. Full comparisons in Table 3.


Quick aside

**CLIP: Contrastive
Language-Image
Pretraining**

(1) Contrastive pre-training



(2) Create dataset classifier from label text



(3) Use for zero-shot prediction

Figure 1. Summary of our approach. While standard image models jointly train an image feature extractor and a linear classifier to predict some label, CLIP jointly trains an image encoder and a text encoder to predict the correct pairings of a batch of (image, text) training examples. At test time the learned text encoder synthesizes a zero-shot linear classifier by embedding the names or descriptions of the target dataset’s classes.

Learning Transferable Visual Models From Natural Language S

```
# image_encoder - ResNet or Vision Transformer
# text_encoder  - CBOW or Text Transformer
# I[n, h, w, c] - minibatch of aligned images
# T[n, l]       - minibatch of aligned texts
# W_i[d_i, d_e] - learned proj of image to embed
# W_t[d_t, d_e] - learned proj of text to embed
# t             - learned temperature parameter

# extract feature representations of each modality
I_f = image_encoder(I) #[n, d_i]
T_f = text_encoder(T)  #[n, d_t]


# joint multimodal embedding [n, d_e]
I_e = l2_normalize(np.dot(I_f, W_i), axis=1)
T_e = l2_normalize(np.dot(T_f, W_t), axis=1)

# scaled pairwise cosine similarities [n, n]
logits = np.dot(I_e, T_e.T) * np.exp(t)

# symmetric loss function
labels = np.arange(n)
loss_i = cross_entropy_loss(logits, labels, axis=0)
loss_t = cross_entropy_loss(logits, labels, axis=1)
loss   = (loss_i + loss_t)/2
```

Figure 3. Numpy-like pseudocode for the core of an implementation of CLIP.

DINO (v1-v ∞)



Previous
SOTA

DINO v1: Self-Distillation w **NO** Labels

Vision transformer backbone – attention learns about objects without labels!





Figure 1: **Self-attention from a Vision Transformer with 8×8 patches trained with no supervision.** We look at the self-attention of the [CLS] token on the heads of the last layer. This token is not attached to any label nor supervision. These maps show that the model automatically learns class-specific features leading to unsupervised object segmentations.

Predict same “label” distribution

(even though we have no g.t. labels)



Augmented views just like before


Mostly same augmentation as SimCLR




(a) Original




(b) Crop and resize



(c) Crop, resize (and flip)



(d) Color distort. (drop)



(e) Color distort. (jitter)



(f) Rotate $\{90^\circ, 180^\circ, 270^\circ\}$



(g) Cutout



(h) Gaussian noise



(i) Gaussian blur




(j) Sobel filtering

+ "solarization"



“self-
distillation”?
Teacher =
Exponential
Moving
Average of
Student?

Predict same “label” distribution
(even though we have no g.t. labels)




Augmented views just like before

v2-3: mostly scaling improvements

Improve the per-pixel or "dense" features




Input



SigLIP 2



PE Spatial



DINOv2 w/reg




DINOv3

Adding patch-level losses

- Student-teacher loss on Patch Features – “Gram Anchoring”
- And iBOT

iBOT (prev.): Self-Distillation on Masked Patches



Problems with vision and uncertainty

Lead in to discussing *uncertainty*



$p(y=\text{"wolf"})=99\%$



$p(y=\text{"lion"})=99\%$





Image classification benchmarks were easy and artificial


Information used by neural nets and humans is not the same!

Hard problems remain.



Sensitivity to adversarial perturbations





x-space (image pixels)

Real world attacks



<https://arxiv.org/pdf/1707.08945.pdf>



Another issue – far Out Of Distribution (OOD) inputs





$p(y=\text{"lion"})=99\%$



$p(y=\text{"lion"})=99\%$








Predictions not invariant to “nuisance” variations




$\hat{y}$ ="lion"



$\hat{y}$ ="wolf"




Parts of image manifold not present in training data can be misclassified




We attempt to mitigate some of these issues with semi- and self-supervised learning

What do the probabilities that our neural nets give us really *mean*?




$p(y=\text{"lion"})=99\%$



$p(y=\text{"lion"})=99\%$



$p(y=\text{"lion"})=99\%$



$p(y=\text{"cat"})=80\%$

We often say that the probability a neural net assigns to the top choice is its **“confidence”**. If a neural net is 80% confident of a prediction, will that prediction be correct 80% of the time?

Confidence calibration game: give a range that you think the answer lies in with probability 50%

Martin Luther King's age at death:

Length of the Nile River:

Number of countries that are members of OPEC:

Number of books in the Old Testament:

My guesses

Correct answers

Martin Luther King's age at death:

35-55

39

Length of the Nile River:

700-1500 miles

6737 km or 4187 miles

Number of countries that are members of OPEC:

9-23


12 (as of 2026)

Number of books in the Old Testament:


30-55

39

Two types of uncertainty



aleatoric / fundamental / irreducible / data uncertainty



Epistemic / Model uncertainty
- due to our lack of knowledge

Do we care? Lots of applications. Right side will cause problems.