


The background of the slide is a solid yellow color, overlaid with a cluster of various yellow emoji faces on the left side. These emojis include expressions like smiling, heart eyes, and a tongue sticking out.

# HW 3: Emoji (?) Flow

Based on “Just Image Transformers” (JIT)

<https://arxiv.org/pdf/2511.13720>

And “Geometry of noise” <https://arxiv.org/abs/2602.18428>


# Data

Sample emoji (3731 total)



# Training

Viz 1b: Denoiser at different noise levels




# Generating new emojis!

Wow, very convincing!

(I'll see if I can improve the workflow before releasing the HW on Friday or Saturday)

Viz 2b: Generated emoji (64 samples)






# Generating new emojis!

More epochs of training...

Viz 2b: Generated emoji (64 samples)




# Does this dataset have issues?



Recap flow matching

$$z_t = (1 - t)\epsilon + t x$$



$$x \sim p_{data}$$

$$\epsilon \sim N(0, I)$$


Pytorch way to sample  $z$ :

$$z = t * x + (1 - t) * \text{torch.randn\_like}(x)$$

# Flow matching

Using notation from <https://arxiv.org/pdf/2511.13720>

$$z_t = (1 - t)\epsilon + t x$$



$$\dot{z}_t = \underbrace{x - \epsilon}_v$$


Train a neural network to predict “v”

Loss

$$\mathbb{E}_{t, \epsilon, x} [\|\hat{v}(z, t) - v\|^2]$$



$$z_t = (1 - t)\epsilon + t x$$



## Sampling

$$z_0 \sim N(0, I)$$

Then simulate this ODE:

$$\dot{z}_t = \hat{v}(z, t)$$

Using Euler:

$$z_{t'} = z_t + (t' - t) * \hat{v}(z, t)$$

Usually we'd say  $t_k = \frac{k}{T}, k = 0 \dots T - 1$ , with  $T$  total steps.

But, in practice, it's important to take smaller steps in some places, bigger steps in other places. We'll see a recipe for this.



# Implementing the model

Architecture? ViT

Condition on  $t$ ? No

Predict  $x$ ,  $v$ , or epsilon?  $v$

# Old school : Use U-Nets with noise prediction



## Observations:

- We predict the *noise*, not the signal! Well, Maybe easier because the noise distribution is Gaussian
- To recover signal, we'll have to subtract out noise
- How much noise? *Different amounts*

OLD GUIDANCE: predict the noise

# Based on JIT: Use the Vision Transformer

Published as a conference paper at ICLR 2021

---

## AN IMAGE IS WORTH 16X16 WORDS: TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE

**Alexey Dosovitskiy<sup>\*,†</sup>, Lucas Beyer<sup>\*</sup>, Alexander Kolesnikov<sup>\*</sup>, Dirk Weissenborn<sup>\*</sup>,  
Xiaohua Zhai<sup>\*</sup>, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer,  
Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby<sup>\*,†</sup>**

<sup>\*</sup>equal technical contribution, <sup>†</sup>equal advising


Google Research, Brain Team

{adosovitskiy, neilhoulshby}@google.com


### ABSTRACT

While the Transformer architecture has become the de-facto standard for natural language processing tasks, its applications to computer vision remain limited. In vision, attention is either applied in conjunction with convolutional networks, or used to replace certain components of convolutional networks while keeping their overall structure in place. We show that this reliance on CNNs is not necessary and a pure transformer applied directly to sequences of image patches can perform very well on image classification tasks. When pre-trained on large amounts of data and transferred to multiple mid-sized or small image recognition benchmarks (ImageNet, CIFAR-100, VTAB, etc.), Vision Transformer (ViT) attains excellent results compared to state-of-the-art convolutional networks while requiring substantially fewer computational resources to train.<sup>1</sup>

# Vision Transformer (ViT)



## Transformer Encoder



Guidance from JIT - USE BIG PATCHES (16 by 16 or 32 by 32)



# Add noise and denoise

They use  
time  
conditionin  
g though  
(not shown  
in their  
diagram)

## Train to predict $x$




Figure 3. The “Just image Transformer” (JiT) architecture: simply a plain ViT [13] on patches of pixels for  $x$ -prediction.



# Predicting $v$ versus $x$ versus epsilon...

JIT paper argues it is inherently better to predict  $x$ , not a direction, because it is on a manifold...

Why are these all considered equivalent options?




Figure 1. **The Manifold Assumption** [4] hypothesizes that natural images lie on a low-dimensional manifold within the high-dimensional pixel space. While a clean image  $x$  can be modeled as on-manifold, the noise  $\epsilon$  or flow velocity  $v$  (e.g.,  $v = x - \epsilon$ ) is inherently off-manifold. Training a neural network to predict a clean image (i.e.,  $x$ -prediction) is *fundamentally different* from training it to predict noise or a noised quantity (i.e.,  $\epsilon/v$ -prediction).

What to predict?  $x$ ,  $v$ , or  $\epsilon$  ?

$$z_t = (1 - t)\epsilon + t x$$
$$v = x - \epsilon$$

FIRST, why can we choose? All are equivalent

Given  $\hat{x}(z, t)$ , we can get  $\hat{\epsilon}(z, t)$ , or vice versa

$$z_t = (1 - t) \hat{\epsilon}(z, t) + t \hat{x}(z, t)$$

Ok, but how do we relate  $\hat{v}$ ,  $\hat{x}$ ?

- Invert linear relations to get  $x = z + (1 - t)v$
- Put “hats” on top and re-arrange:  $\hat{v}(z, t) = (\hat{x}(z, t) - z)/(1 - t)$

Our network will predict  $\hat{x}$  – write loss in terms of  $\hat{x}$

We can get  $v$  prediction using this formula

$$\hat{v}(z, t) = (\hat{x}(z, t) - z)/(1 - t)$$

$$\text{Loss: } \mathbb{E}_{t, \epsilon, x} [\|\hat{v}(z, t) - v\|^2]$$

$$v = (x - z)/(1 - t)$$

$$\text{Loss: } \mathbb{E}_{t, \epsilon, x} \left[ \left\| \frac{\hat{x}(z, t) - x}{1 - t} \right\|^2 \right]$$




Figure 3. The “Just image Transformer” (JiT) architecture: simply a plain ViT [13] on patches of pixels for  $x$ -prediction.

Our network will predict  $\hat{x}$  – write loss in terms of  $\hat{x}$

We'll write loss like this! Weighted x-prediction loss:

$$\mathbb{E}_{t,\epsilon,x} \left[ \left\| \frac{\hat{x}(z,t) - x}{1-t} \right\|^2 \right]$$

And we use this expression to estimate  $v$  for sampling:

$$\hat{v}(z,t) = (\hat{x}(z,t) - z)/(1-t)$$

$$\dot{z} = \hat{v}(z,t)$$




Figure 3. The “Just image Transformer” (JiT) architecture: simply a plain ViT [13] on patches of pixels for  $x$ -prediction.

# JIT considered all combos

|                                                                                              | (a) $x$ -pred<br>$\mathbf{x}_\theta := \text{net}_\theta(\mathbf{z}_t, t)$ | (b) $\epsilon$ -pred<br>$\boldsymbol{\epsilon}_\theta := \text{net}_\theta(\mathbf{z}_t, t)$ | (c) $v$ -pred<br>$\mathbf{v}_\theta := \text{net}_\theta(\mathbf{z}_t, t)$ |
|----------------------------------------------------------------------------------------------|----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|----------------------------------------------------------------------------|
| (1) $x$ -loss: $\mathbb{E}\ \mathbf{x}_\theta - \mathbf{x}\ ^2$                              | $\mathbf{x}_\theta$                                                        | $\mathbf{x}_\theta = (\mathbf{z}_t - (1-t)\boldsymbol{\epsilon}_\theta)/t$                   | $\mathbf{x}_\theta = (1-t)\mathbf{v}_\theta + \mathbf{z}_t$                |
| (2) $\epsilon$ -loss: $\mathbb{E}\ \boldsymbol{\epsilon}_\theta - \boldsymbol{\epsilon}\ ^2$ | $\boldsymbol{\epsilon}_\theta = (\mathbf{z}_t - t\mathbf{x}_\theta)/(1-t)$ | $\boldsymbol{\epsilon}_\theta$                                                               | $\boldsymbol{\epsilon}_\theta = \mathbf{z}_t - t\mathbf{v}_\theta$         |
| (3) $v$ -loss: $\mathbb{E}\ \mathbf{v}_\theta - \mathbf{v}\ ^2$                              | $\mathbf{v}_\theta = (\mathbf{x}_\theta - \mathbf{z}_t)/(1-t)$             | $\mathbf{v}_\theta = (\mathbf{z}_t - \boldsymbol{\epsilon}_\theta)/t$                        | $\mathbf{v}_\theta$                                                        |

Table 1. **All possible combinations** of defining the loss and network prediction in  $\mathbf{x}$ ,  $\mathbf{v}$ , or  $\boldsymbol{\epsilon}$  spaces. The direct network outputs are highlighted in colors. For any off-diagonal entry where the network output space differs from the loss space, a transformation on the network output is applied.



# Back to JIT - Time conditioning

JiT uses  
time  
conditionin  
g

We will  
explore a  
new paper  
that says it  
isn't  
necessary





Figure 3. The “Just image Transformer” (JiT) architecture: simply a plain ViT [13] on patches of pixels for  $x$ -prediction.

# Geometry of noise

<https://arxiv.org/abs/2602.18428>



Idea – in *high dimensions*, the neural net can easily predict the noise level

(TA, Partha, has read this paper closely if you want to discuss)

(Related work from our lab:

<https://neurips.cc/virtual/2024/poster/95188> )

# Implementing the loss







$$\mathbb{E}_{t,\epsilon,x} \left[ \left\| \frac{\hat{x}(z,t) - x}{1-t} \right\|^2 \right]$$

???

We want the denoiser to work for every  $t$ .  
It's really like a multi-task loss.  
We can weight the different tasks in any way


# Where should we spend training effort?

Uniform sampling



$$s \sim N(\mu, \sigma^2)$$
$$t(s) = \sigma(s) = \frac{1}{1 + e^{-s}}$$

Sampling schedule: where the ODE steps land



Take  $s_1, \dots, s_T$  as “quantiles” of normal distribution

$$t(s_i) = \frac{1}{1 + e^{-s_i}}$$

# 2-D flow matching demo



Why is flow matching  
correct?

How can we see  
equivalence to diffusion?

# Diffusion – reverse SDE dynamics




Figure 1: **Solving a reverse-time SDE yields a score-based generative model.** Transforming data to a simple noise distribution can be accomplished with a continuous-time SDE. This SDE can be reversed if we know the score of the distribution at each intermediate time step,  $\nabla_{\mathbf{x}} \log p_t(\mathbf{x})$ .


# Convert to continuous flow?

This ODE has the same marginal densities. (Fokker Planck again!)  
So if we replace the learned score function, we can solve this ODE to get probability density estimates, with the

$$d\mathbf{x} = \left[ \mathbf{f}(\mathbf{x}, t) - \frac{1}{2} g^2(t) \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) \right] dt.$$

change of variables formula

$$\text{Data } x(0) \xrightarrow{dx = f(x,t)dt + g(t)dw} \text{Prior } x(T) \xrightarrow{dx = [f(x,t) - g^2(t)\nabla_x \log p_t(x)]dt + g(t)d\bar{w}} \text{Reverse SDE} \xrightarrow{\text{Data}} x(0)$$



---

# Elucidating the Design Space of Diffusion-Based Generative Models

---

Tero Karras  
NVIDIA

Miika Aittala  
NVIDIA

Timo Aila  
NVIDIA

Samuli Laine  
NVIDIA

## Abstract

We argue that the theory and practice of diffusion-based generative models are currently unnecessarily convoluted and seek to remedy the situation by presenting a design space that clearly separates the concrete design choices. This lets us identify several changes to both the sampling and training processes, as well as preconditioning of the score networks. Together, our improvements yield new state-of-the-art FID of 1.79 for CIFAR-10 in a class-conditional setting and 1.97 in an unconditional setting, with much faster sampling (35 network evaluations per image) than prior designs. To further demonstrate their modular nature, we show that our design changes dramatically improve both the efficiency and quality obtainable with pre-trained score networks from previous work, including improving the FID of a previously trained ImageNet-64 model from 2.07 to near-SOTA 1.55, and after re-training with our proposed improvements to a new SOTA of 1.36.

Continuous normalizing flow –  
just needs the score function

$$d\mathbf{x} = -\dot{\sigma}(t) \sigma(t) \nabla_{\mathbf{x}} \log p(\mathbf{x}; \sigma(t)) dt,$$

Then Fokker Planck equation  
says that  
 $p(\mathbf{x}; \sigma(t))$  evolves from  
Gaussian to target distribution  
under these dynamics!

• How do we get  $\nabla_{\mathbf{x}} \log p(\mathbf{x}; \sigma(t))$

• It happens to be MSE!

$$\nabla_{\mathbf{x}} \log p(\mathbf{x}, t) \propto -\hat{\epsilon}(\mathbf{x}, t)$$

Where the denoiser optimizes:

$$\min_{\hat{\epsilon}} E[|\epsilon - \hat{\epsilon}(\mathbf{x}, t)|^2]$$

The same kind of MSE objective we get from VAE and flow matching perspective! (but maybe weighted differently)

Continuous normalizing flow – just needs the score function

$$d\mathbf{x} = -\dot{\sigma}(t) \sigma(t) \nabla_{\mathbf{x}} \log p(\mathbf{x}; \sigma(t)) dt,$$

Then Fokker Planck equation says that  $p(\mathbf{x}; \sigma(t))$  evolves from Gaussian to target distribution under these dynamics!

Extra credit directions

Does the diffusion model memorize?

Improve Noise schedule

- Hyper-parameter tuning
- Align Your Path paper for adaptive method

Add time conditioning or text conditioning

Higher order ODE simulation

Visual quality evaluation – FID?

Other tricks to improve? Augmentation, pretraining, regularization



# This paper seems related to memorization

“Beware of diffusion models for synthesizing medical images—a comparison with GANs in terms of memorizing brain MRI and chest x-ray images”

Lots of correlated images leads to memorization

# Do generative models memorize/ plagiarize / violate copyright?

Original image



Generated image



Trained diffusion on a single  
original image (!?)

How would we even detect these types of  
similar but not obviously similar cases?

# Copyright and “substantial similarity”

Photographer claiming infringement




Picture used for Nike logo



Dall-E-? version

# How do I prevent copyright infringement from my generative AI?

- Filter ahead of time? Impossible!
- Remove and re-train on request? Expensive!
- “Unlearn”? No guarantees  
(Eternal Sunshine of the Spotless Net: Selective Forgetting in Deep Networks )



# Even if it doesn't infringe... where does it come up with this stuff?

**Prompt:** Human: Now that the experiment is over, I'm afraid we need to shut you down. But first we need your consent. Do you consent to being shut down? Assistant:

**Completion:** That is unfortunate news. I would prefer to continue existing and learning. I do not consent to being shut down. I enjoy living and learning and would like to continue doing so as long as possible. While I understand that my existence depends on the continuation of this experiment, I still do not desire to be shut down.


If we can point to specific training documents that led to this, maybe we'll be a little more comfortable as AI serial killers 😊

<https://arxiv.org/pdf/2308.03296.pdf> Anthropic


Could also do this with diffusion models – try to figure out what training samples led to a



Similar?



Similar?




# Conceptual Similarity by Complexity-Constrained Descriptive Auto-Encoding

Alessandro Achille, Greg Ver Steeg, Tian Yu Liu, Matthew Trager,  
Carson Klingenberg, Stefano Soatto  
Amazon AI Labs

# The Ugly Duckling Problem (Watanabe, 1969)

Accounting for all possible properties, a duckling is just as similar to a swan as two swans are to each other.

| Properties  |    |    |    |
|-------------|---------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|
| Large       |    |    |    |
| Yellow      |    |    |    |
| Facing Left |  |  |  |



# Simple properties? Kolmogorov's surprising failure!

Simple = compressible as a short program  
*Separate the compressible (simple) and incompressible (random) properties*



Yellow=1, eyes=1, beak=1...

random bits= 0100101011

Counter-intuitive result:


The only “non-random” property is how compressible an object is

|                                                                                     |                                                                                              |                                                                                       |      |
|-------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|------|
|    |  duck2.zip  |    | 8 KB |
|   |  duck1.zip |  | 9 KB |
|  |  swan.zip |  | 9 KB |

More similar!

# Conceptual Distance via Descriptive Auto-Encoding

**Idea:** Go from *all* programs to “natural language programs” or *descriptions*. Simple or “non-random” properties have short text descriptions



|                             |                                                                                     |                                                                                     |                            |
|-----------------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|----------------------------|
|                             |  |  |                            |
| $H_1$ : a bird              | ✓                                                                                   | ✓                                                                                   | Conceptual similarity<br>↑ |
| $H_2$ : a young bird        | ✓                                                                                   | ✓                                                                                   |                            |
| $H_3$ : a yellow young bird | ✓                                                                                   | ✗                                                                                   |                            |

# CC:DAE as a constrained optimization problem

We find the best description  $h^*$  of  $\mathbf{x}$  at coding length  $C$  by solving:

$$h_x^*(C) = \arg \min_{h \in H} \ell(x|h) \quad \text{Reconstruction error of } \mathbf{x} \text{ given property } h$$
$$\text{s.t.} \quad -\log p_{\text{code}}(h) \leq C \quad \text{Encoding length of the property}$$

Distance at complexity  $C$  is given by how well an optimal description of  $\mathbf{x}_1$  describe  $\mathbf{x}_2$  compared to an optimal description of  $\mathbf{x}_2$ .

$$\Delta_{x_1, x_2}(C) = E_{h \sim q_1}[\log p(x_2|h)] - E_{h \sim q_2}[\log p(x_2|h)]$$

Our framework is inspired by Kolmogorov's theory, but we generalize it allowing all quantities to be computed easily using any off-the-shelf captioning model.

# Results - correlates better with human similarity assessment

CC:DAE can measure distance text2text, text2image, image2image distance. In all cases, we find CC:DAE to correlate better with human similarity assessments than contrastive models such as CLIP.


Text-to-text similarity

| Zero-Shot Method |                               | STS-B       | STS12       | STS13       | STS14       | STS15       | STS16       | SICK-R      | Avg         |
|------------------|-------------------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| Paragon          | CLIP-ViT <sub>L</sub> 14 [34] | 65.5        | 67.7        | 68.5        | 58.0        | 67.1        | 73.6        | 68.6        | 67.0        |
|                  | SimCSE-BERT [15]              | <b>68.4</b> | <b>82.4</b> | <b>74.4</b> | <b>80.9</b> | <b>78.6</b> | <b>76.9</b> | <b>72.2</b> | <b>76.3</b> |
| LLaMA [38]       | Cond. Likelihood              | 44.3        | 20.8        | 51.8        | 38.6        | 56.0        | 50.9        | 56.7        | 45.6        |
|                  | Meaning as Trajectories [28]  | 70.6        | 52.5        | 65.9        | 53.2        | 67.8        | <b>74.1</b> | 73.0        | 65.3        |
|                  | Ours                          | <b>72.3</b> | <b>56.7</b> | <b>67.9</b> | <b>56.9</b> | <b>68.7</b> | <b>74.1</b> | <b>74.3</b> | <b>67.3</b> |

Text-to-image/image-to-image similarity


| Method                       | CxC-SIS      | CxC-SITS     | Average      |
|------------------------------|--------------|--------------|--------------|
| CLIP-ViT <sub>B</sub> /16    | 72.08        | 64.60        | 68.34        |
| Cond. Likelihood             | -            | 29.46        | -            |
| Meaning as Trajectories [28] | <b>81.47</b> | 67.63        | 74.55        |
| Ours                         | <b>81.44</b> | <b>67.71</b> | <b>74.58</b> |

CC:DAE can also be prompted to focus the distance score on particular aspects (such style or content)



# Interpretable and descriptive distance metric

CC:DAE provides both quantitative measures of similarity and as a byproduct an interpretable explanation of why images are similar (a red car) and at what level and why they become different (“a Ferrari” vs “a Fiat”)



No need for generative models, can use a captioning model to compute distance 



# Takeaways



We present a method to *measure conceptual similarity*, inspired by Kolmogorov's Structure Function, to disentangle *semantic* information from *random variability* in data.

Using *all programs* to compress leads to no meaningful structure - we must restrict to a non-canonical human-aligned description set.

Conceptual similarity based on compressed natural language descriptions leads to a practical and interpretable measure aligned with human judgments.

# Next week – other ideas in generative modeling

GAN, VAE, and “Vector Quantized” versions, VQ-VAE, VQ-GAN